

# Monitoring

Fish  
jfsc@cin.ufpe.br

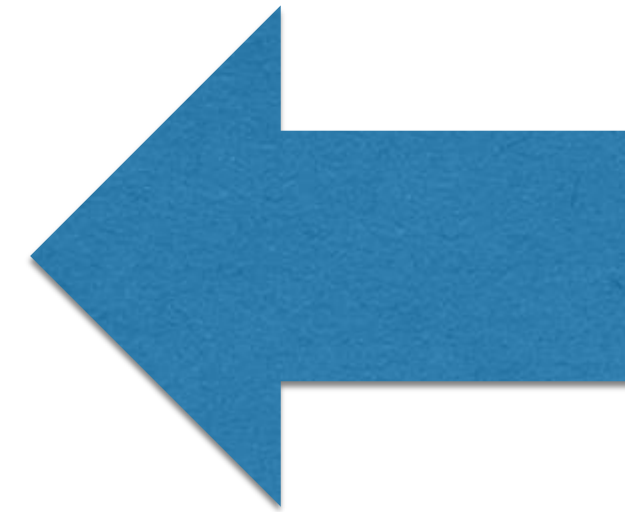
# THE ATTENDANTS WILL LEARN...

- Monitoring basics
- Traditional vs Modern monitoring
- Challenges of monitoring at scale
- What to monitor for DevOps?
- Microservices monitoring rules
- Elastic Stack fundamentals

# What is monitoring?

# Monitoring Topics

- Customer happiness
- Cost efficiency
- Safety and security
- Compliance



We'll talk about this for this class

# Customer Happiness

- Time to value (features in front of client)
- Availability
- Response time

# Cost Efficiency

- Utilization
- Optimization
- Automation



# Traditional Monitoring

- The traditional monitoring approach tests just whether something is 'up' or 'down';
- Agents inside: Hardware, Mainframe, VMs



# Modern monitoring

- If a single component within a microservice architecture fails, there may be no **business impact**, and so **the severity of alert should match this fact**.
- Containers, events (serverless).

# Challenges: Aggregation

- Consolidating different (or without) log formats coming from a lot of containers or events with different proposal of visibility

# Distribution: Interdependence

- Maintain operational health and doing Root Cause Analysis on a large ecosystem (many containers or events) at scale becomes challenging.

# Frequency of change: ephemerality

---

- The short lifetime of containers or events has been considered a big challenge in modern monitoring.

# Unpredictable

- We don't know where the container. Has the Kubernetes moved Nginx to another zone?

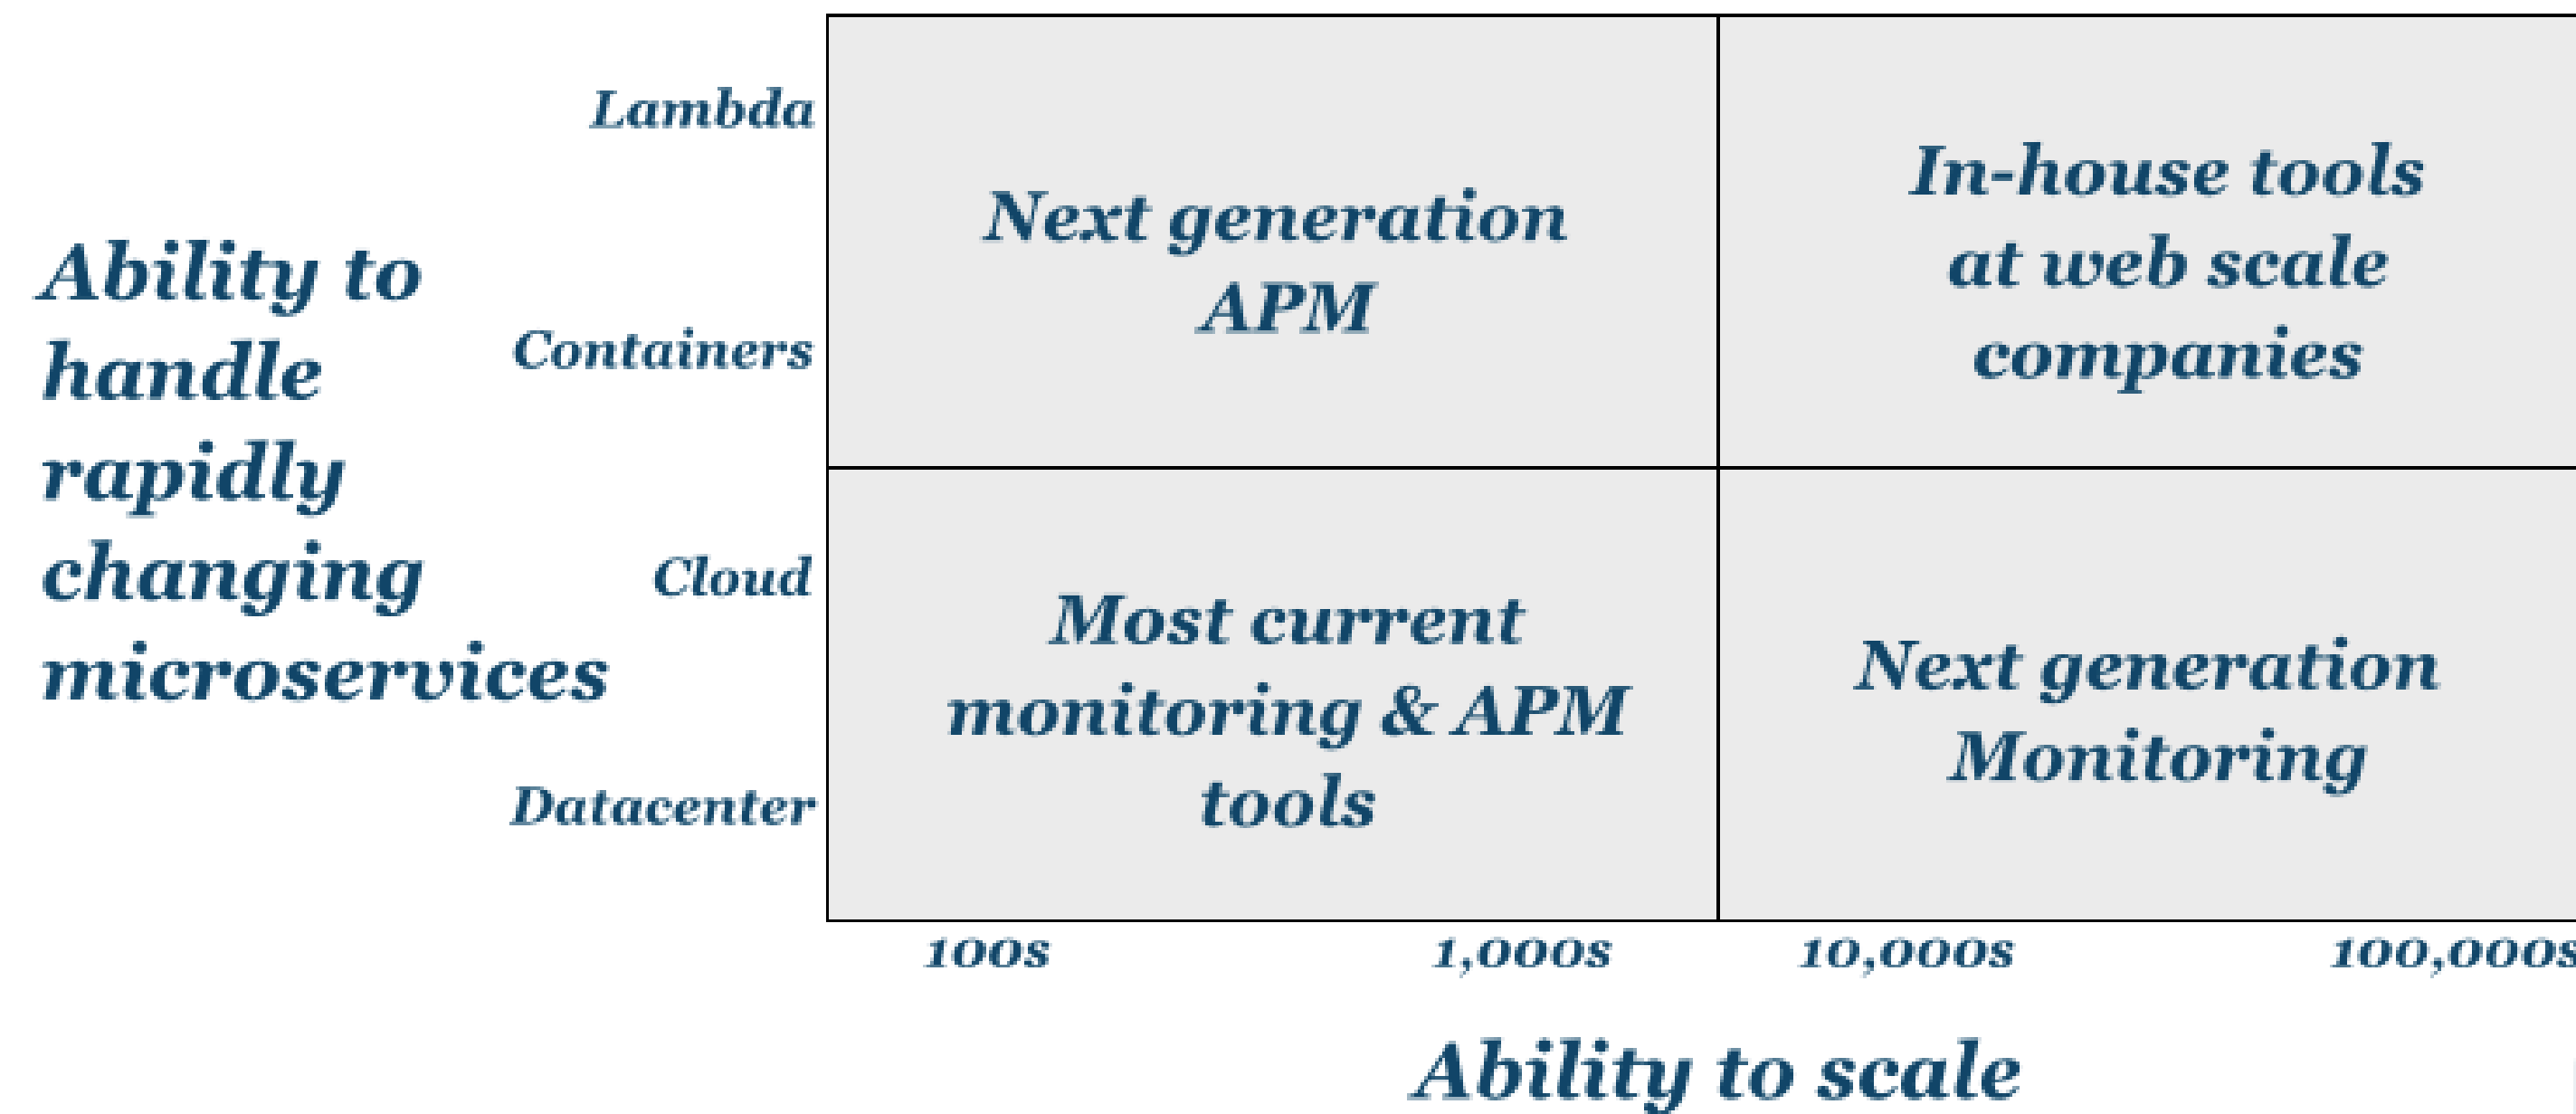


# Tips: 7 containers per host



# The Adrian's Tragic Quadrant

## A Tragic Quadrant



From: <https://www.battery.com/powered/bigpanda-analyzing-the-monitoringscape/>

# What to monitor? from orquestrators

- Swarm,
- mesos,
- kubernetes



# What to monitor? anomalies

---

- Monitoring systems have to learn normal behavior and indicate when system behavior is not normal.

# Rules for microservices

- Spend more time working on code that analyses the meaning of metrics than code that collects, moves, stores and displays metrics
- Reduce key business metric latency to less than the human attention span (~10s)
- Validate your measurement system has enough accuracy and precision. Collect histograms of response time
- Monitoring systems need to be more available and scalable than the systems (and services) being monitored
- Optimise for monitoring distributed, ephemeral, 'cloud-native', containerised microservices
- Fit metrics to models in order to understand relationships (this is a new rule)

From: <https://github.com/adrianco>

# Hands on phase

# Monitoring with Elastic Stack

# What is a log?

# A log is an event

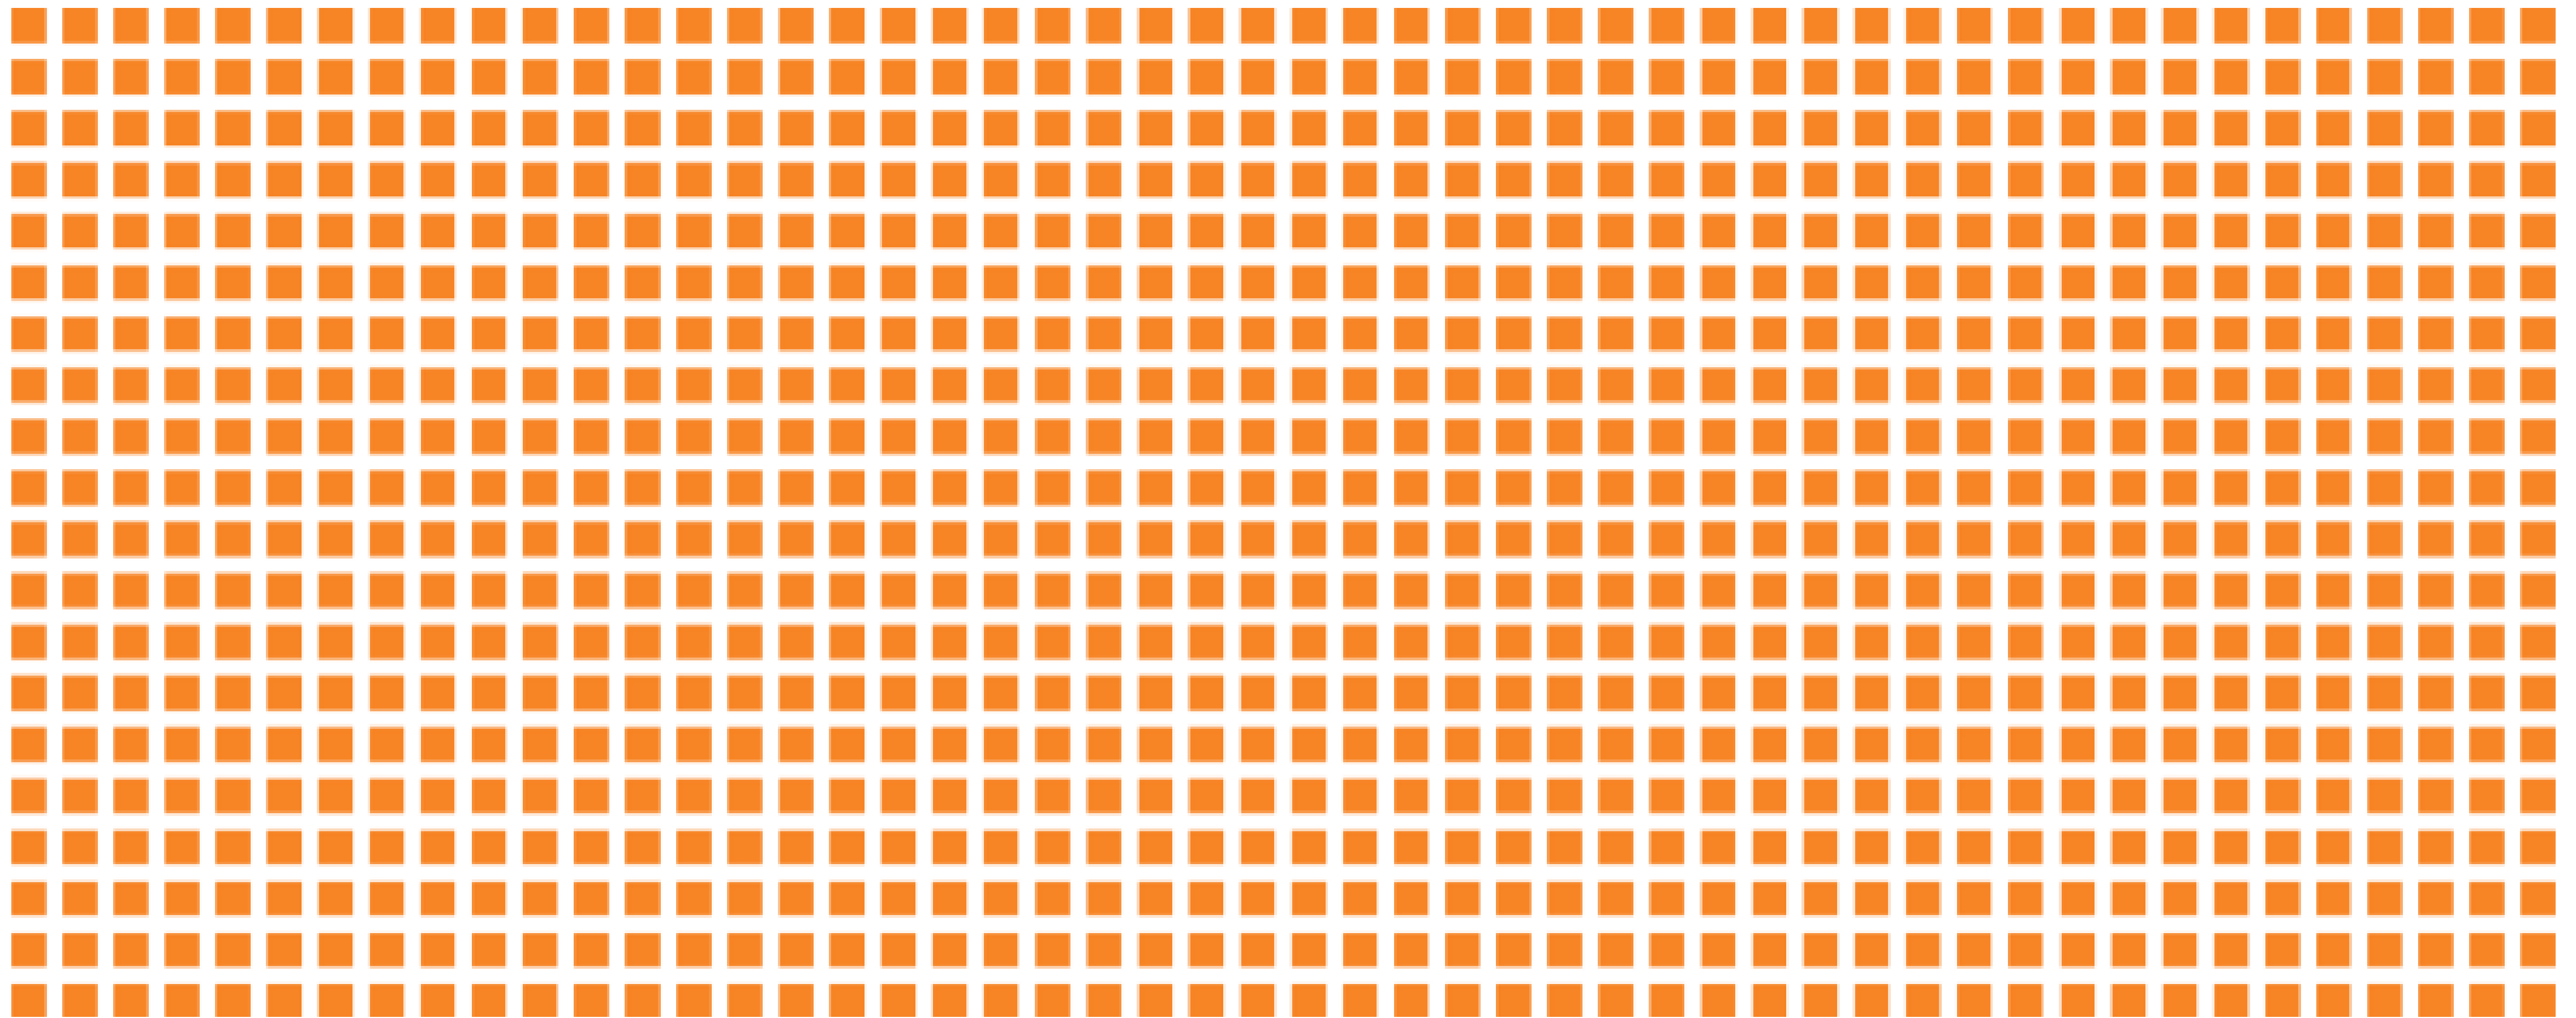
- Log lines
- Twitter feed
- Invoices
- Metrics

# Exercise 1

---

- Reading logs from containers

# At scale



1000 EC2 Instances in a Cluster



# grep

- Multiple machines
- Multiple logs
- Analysis/Discovery
- Time Period

# Different point of views

apache

```
[23/Jan/2014:17:11:55 +0000]
```

unix timestamp

```
1390994740
```

log4j

```
[2014-01-29 12:28:25,470]
```

postfix.log

```
Feb 3 20:37:35
```

ISO 8601

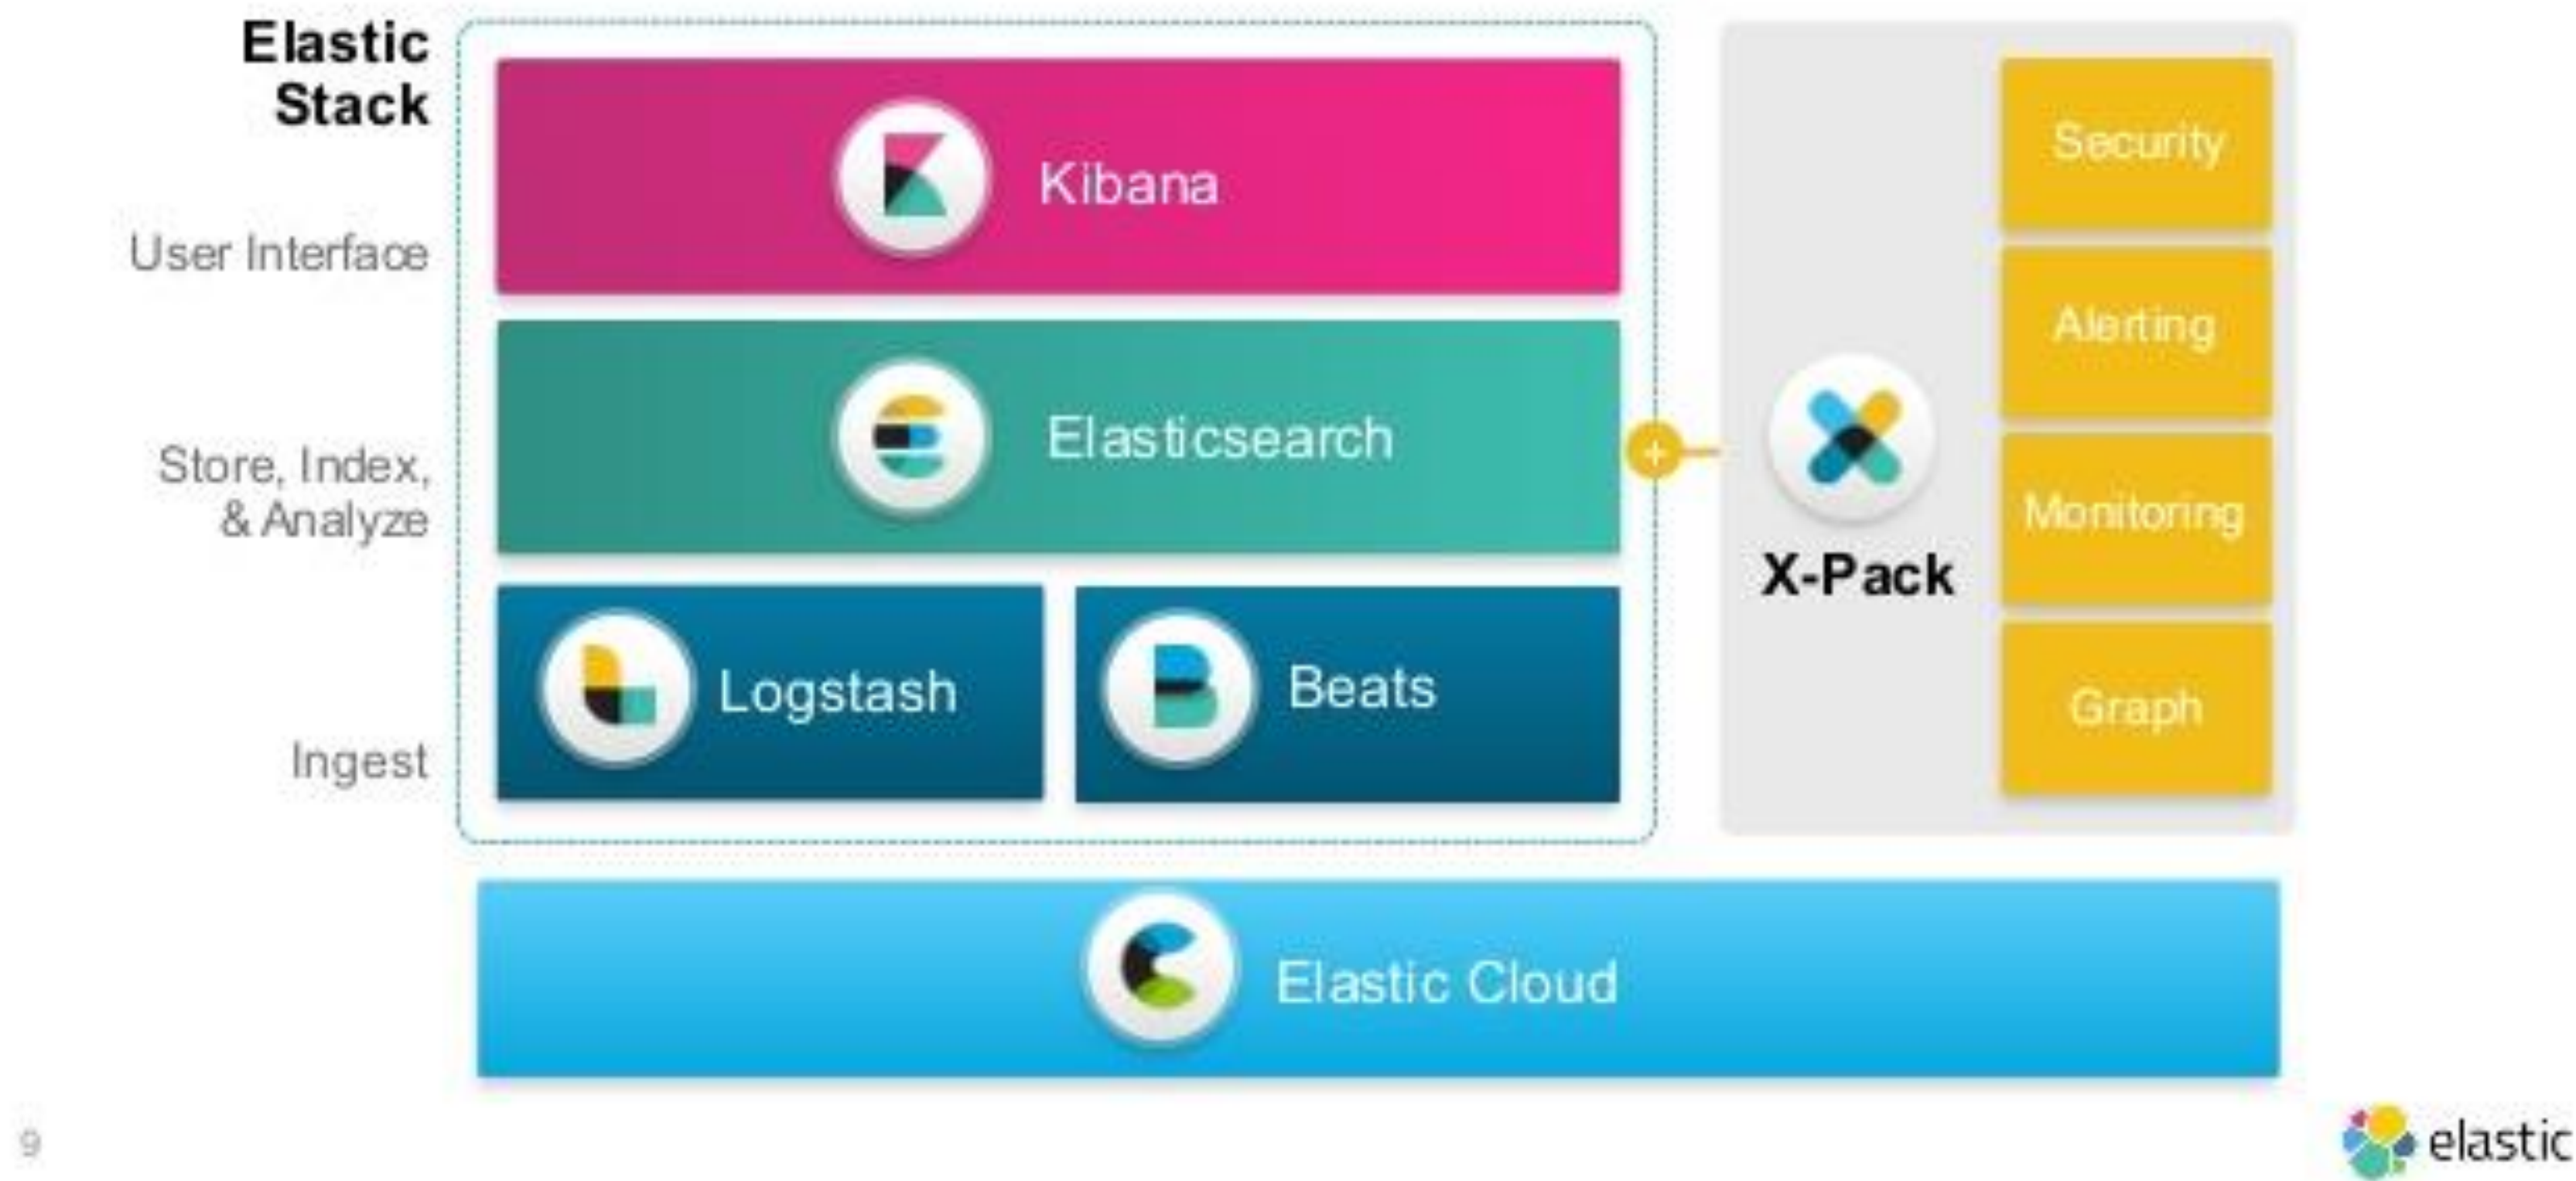
```
2009-01-01T12:00:00+01:00
```

# A Centralized logging approach can help you

- Collect data
- Parse data
- Enrich data
- Store data
- Search and aggregate
- Visualize data

# An ecosystem as response: Elastic Stack

The Elastic Stack is a set of tools developed by Elastic to provide monitoring (nowadays).



# Some Elastic Stack features

- A business centric monitoring;
- Integration-driven (flexibility to get logs from everywhere);
- Taming decentralization;
- Process (indexing) and search near real-time data;
- Use analytics to summarize data across dimensions



# Logstash : inputs

<b>Monitoring</b>	collectd, graphite, ganglia, snmptrap, zenoss
<b>Datastores</b>	elasticsearch, redis, sqlite, s3
<b>Queues</b>	kafka, rabbitmq, zeromq
<b>Logging</b>	beats, eventlog, gelf, log4j, rlp, syslog, varnish log
<b>Platforms</b>	drupal_dblog, gemfire, heroku, sqs, s3, twitter
<b>Local</b>	exec, generator, file, stdin, pipe, unix
<b>Protocol</b>	imap, irc, stomp, tcp, udp, websocket, wmi, xmpp

# Logstash : Filters

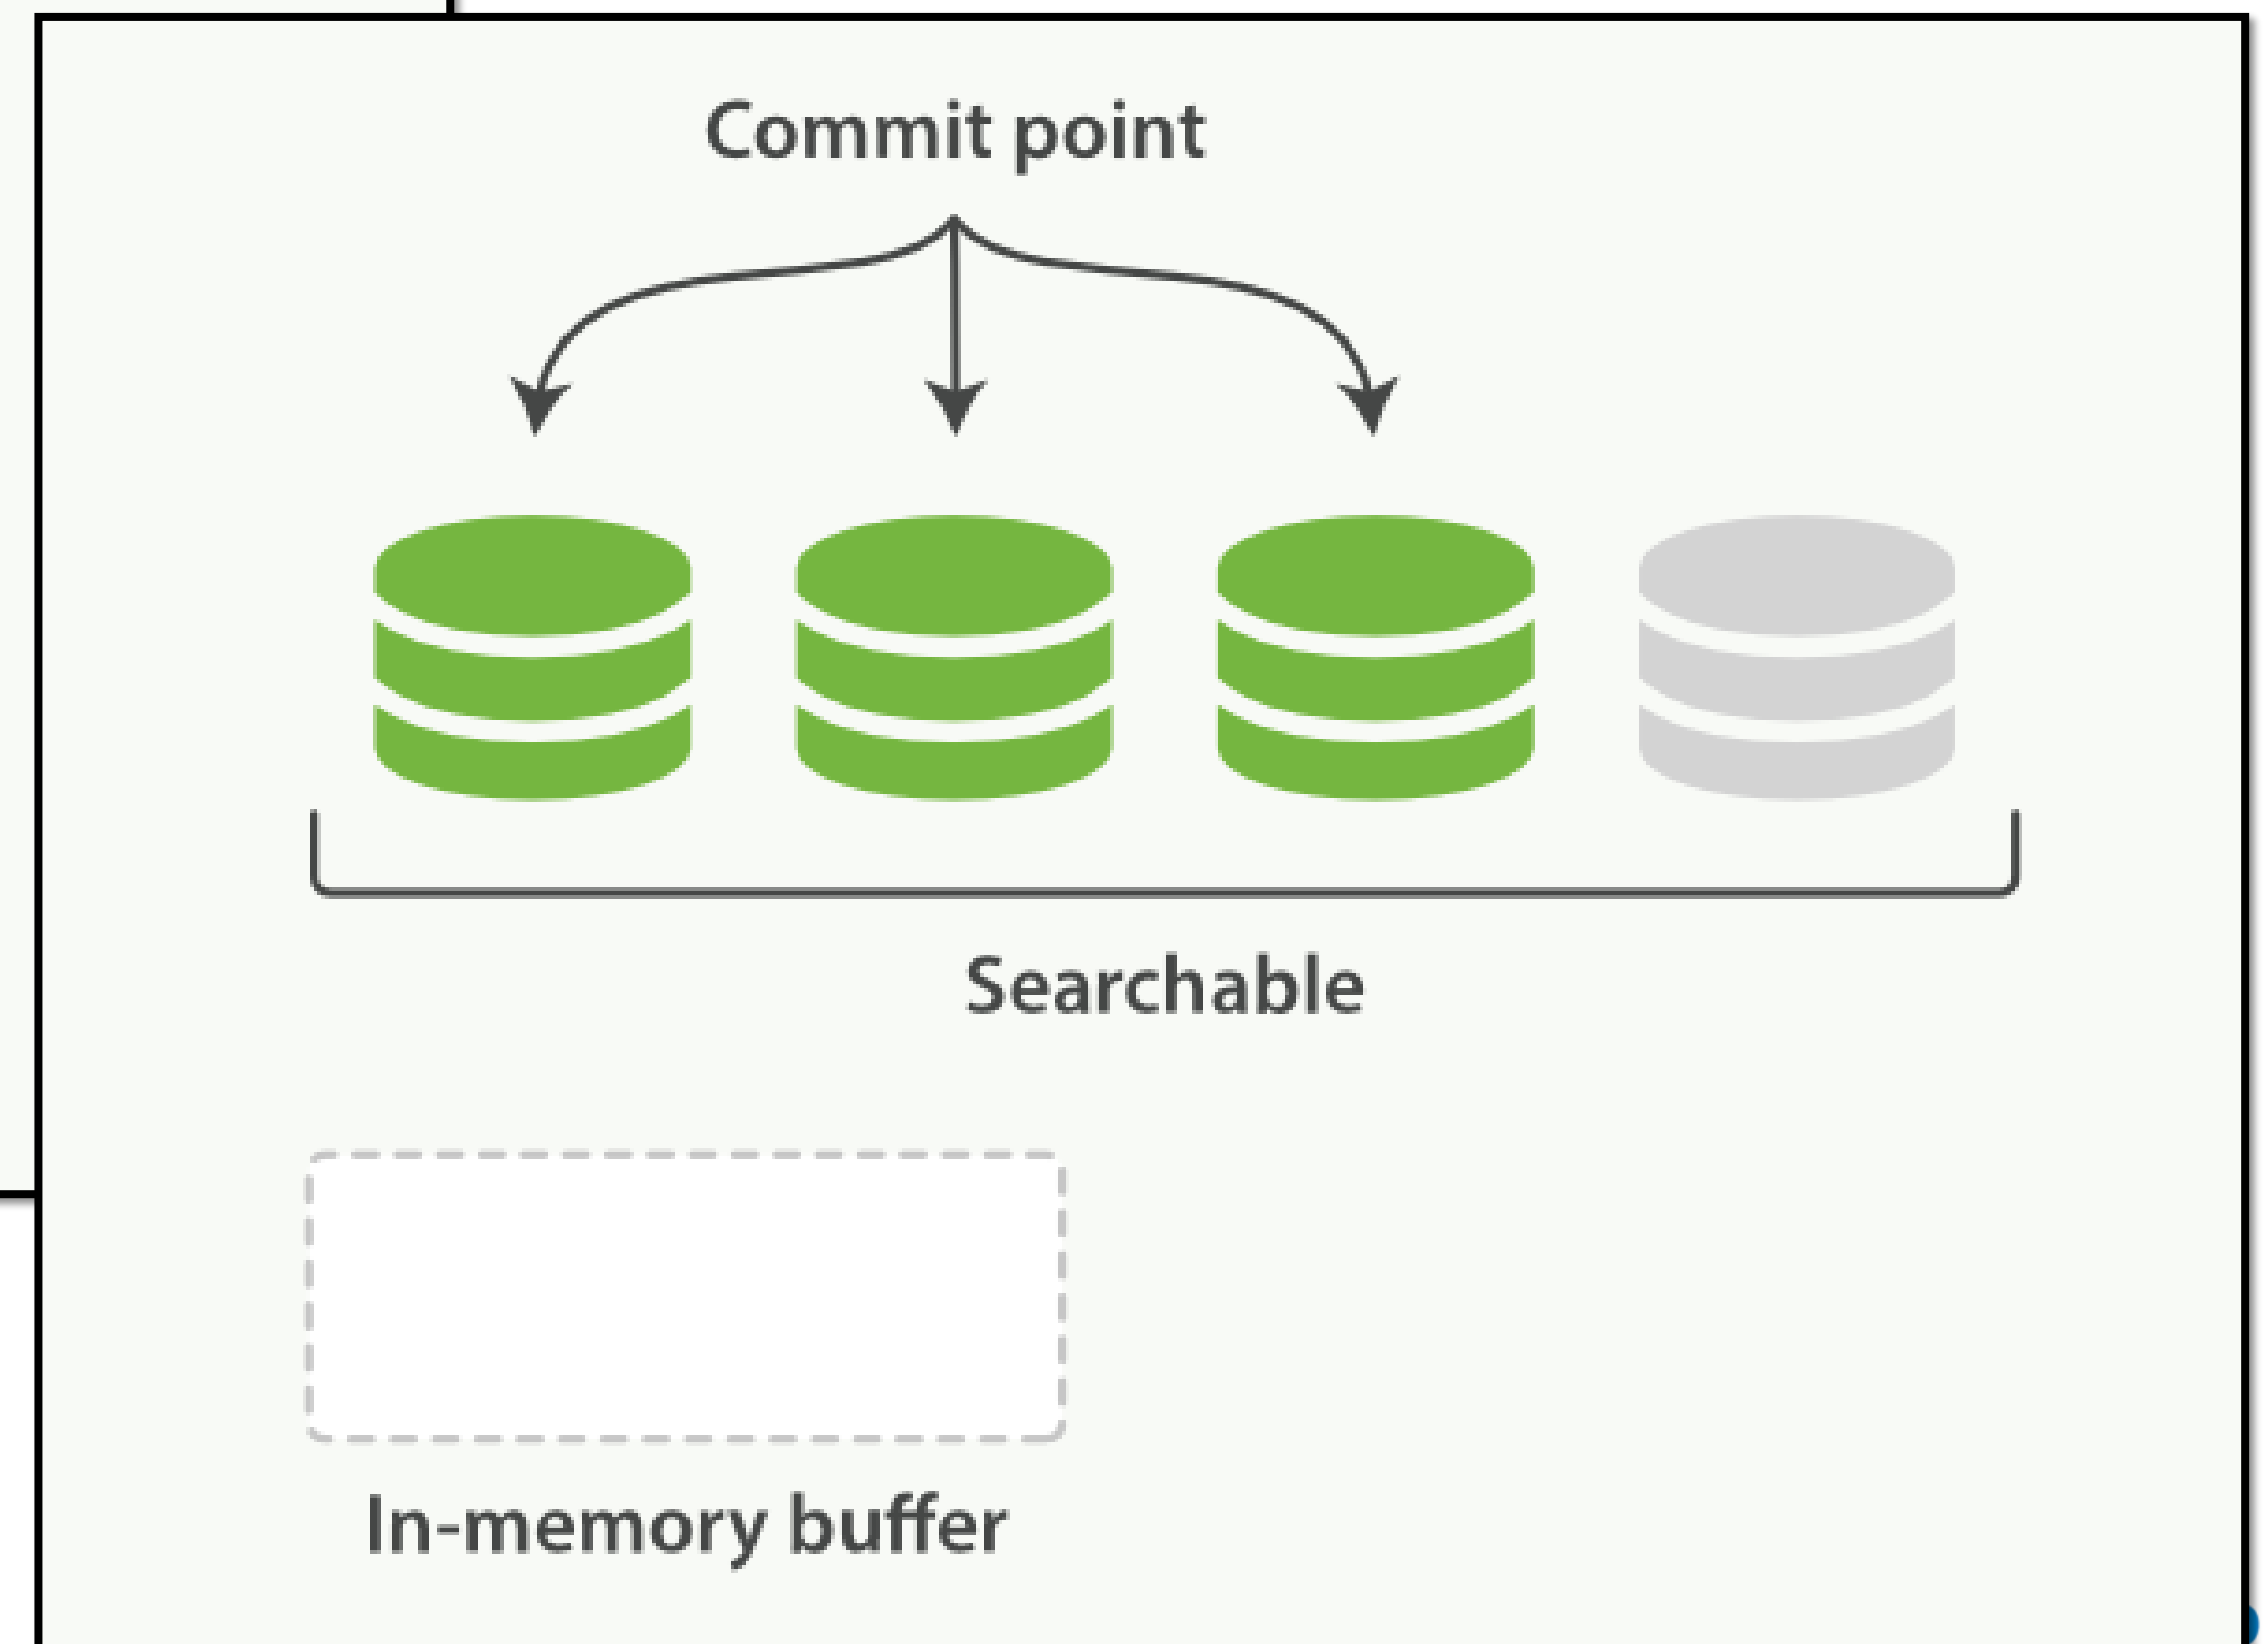
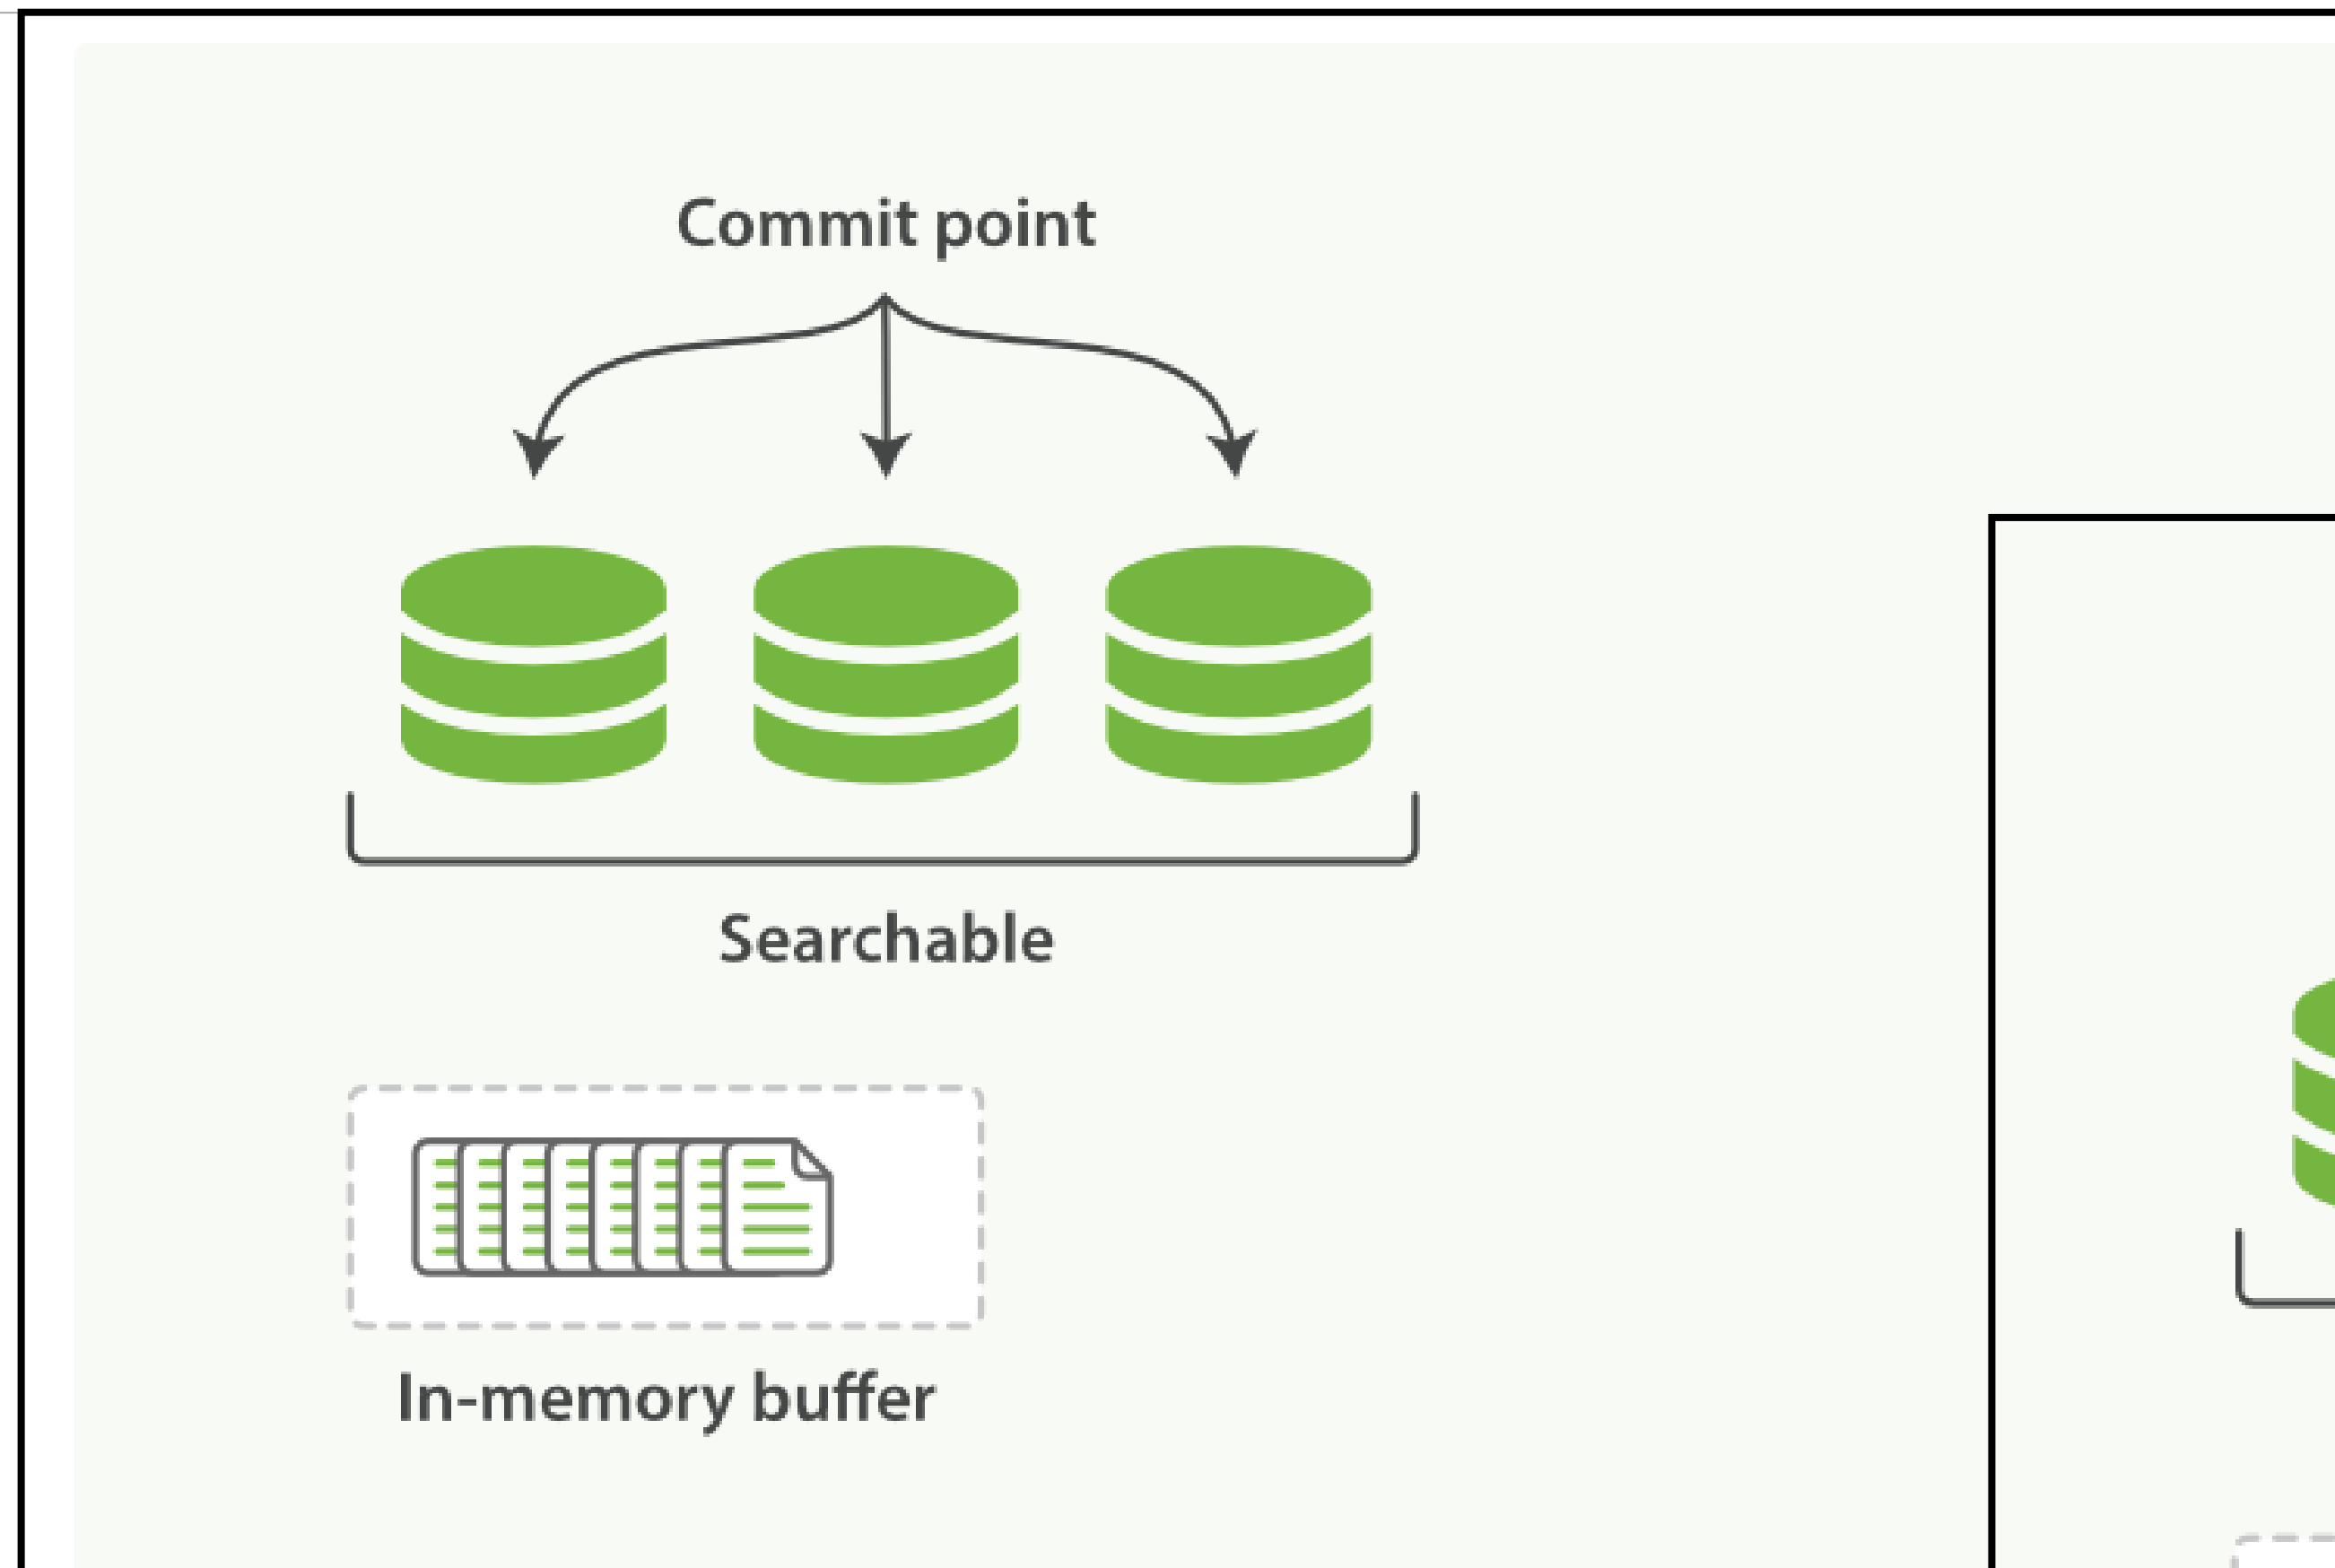
aggregate alter anonymize collate csv  
cidr clone cipher checksum date dns  
drop elasticsearch extractnumbers  
environment elapsed fingerprint geoip  
grok i18n json json\_encode kv mutate  
metrics multiline metaevent prune punct  
ruby range syslog\_pri sleep split throttle  
translate uuid urldecode useragent xml

# Logstash : Output

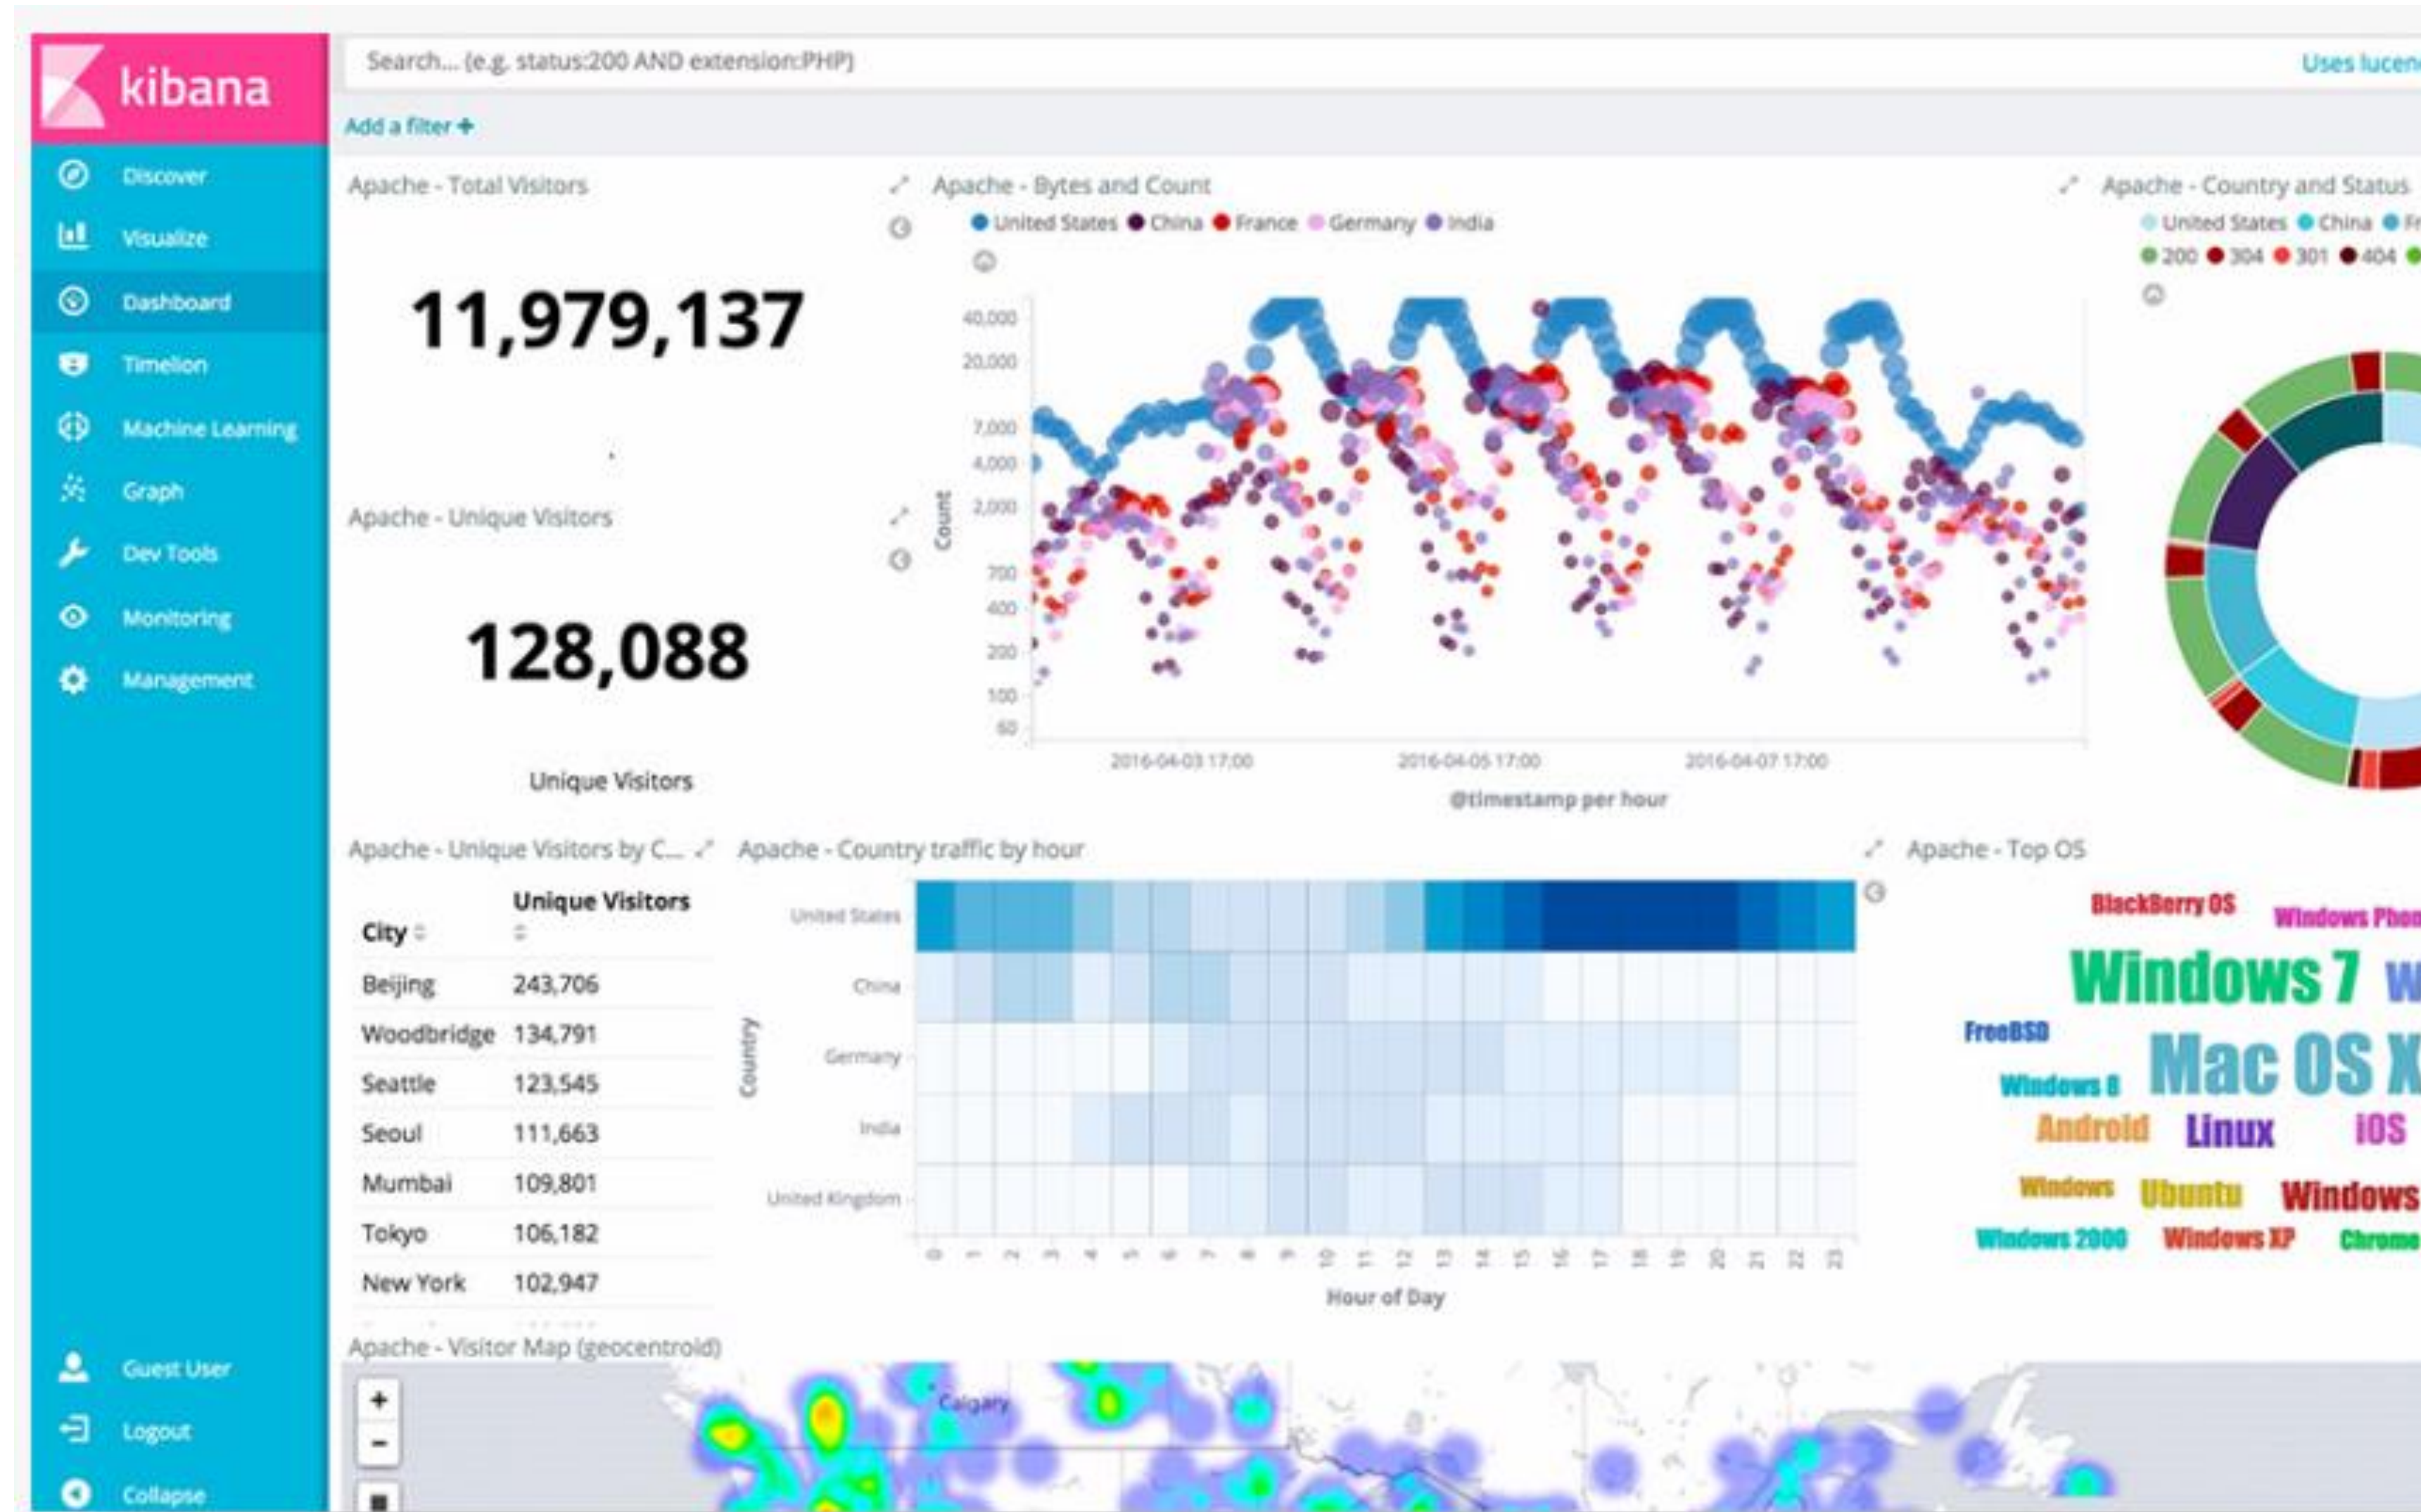
<b>Store</b>	elasticsearch, gemfire, mongodb, redis, riak, rabbitmq, solr
<b>Monitoring</b>	ganglia, graphite, graphtastic, nagios, opentsdb, statsd, zabbix
<b>Notification</b>	email, hipchat, irc, pagerduty, sns
<b>Protocol</b>	gelf, http, lumberjack, metriccatcher, stomp, tcp, udp, websocket, xmpp
<b>External service</b>	google big query, google cloud storage, jira, loggly, riemann, s3, sqs, syslog, datadog
<b>External monitoring</b>	boundary, circonus, cloudwatch, librato
<b>Local</b>	csv, dots, exec, file, pipe, stdout, null



# Elastic Search



# Kibana





# References

- <https://vimeo.com/173609948>
- <https://pt.slideshare.net/adriancockcroft/monitoring-challenges-monitorama-2016-monitoringless>
- <https://speakerdeck.com/elasticsearch/log-all-the-things>
- <https://www.battery.com/powered/bigpanda-analyzing-the-monitoringscape/>